



# CKAD Crash Course

Sander van Vugt  
[mail@sandervanvugt.nl](mailto:mail@sandervanvugt.nl)

# About Me

- Sander van Vugt
- Living in the Netherlands
- Author and presenter of many titles on this platform – Linux, Kubernetes and Ansible
- Founder of the Living Open Source Foundation
  - The mission of the Living Open Source Foundation is to stimulate the growth of local economies by enabling people to develop themselves as experts in the area of Open Source
  - Current focus is on education in Africa
  - See [livingopensource.net](http://livingopensource.net) for more information

# Course resources

- YAML files used in this course are available at  
<https://github.com/sandervanvugt/ckad>

# Poll Question 1

How would you rate your current Kubernetes Knowledge?

- None
- Beginner
- Intermediate
- Advanced
- Guru

# Poll Question 2

How would you rate your current experience with containers?

- None
- Beginner
- Intermediate
- Advanced
- Guru

# Poll Question 3

Which of the following statements best applies to you

- I'm new to Kubernetes, teach me everything you know!
- Give me an overview of Kubernetes, I want to understand what it is doing
- Give me hands-on! I want to learn how it works
- Focus on the exam questions! I already know my basics

# Agenda

- Creating the Lab Environment
- Understanding Kubernetes
- Using Management Options
- Managing Pods
- Managing Deployments
- Managing Services
- Managing Ingress
- Managing Storage
- Managing ConfigMaps and Secrets
- Managing ServiceAccount
- Troubleshooting

# This course vs my other Kubernetes courses

- Managing Containers on Linux: provides knowledge about the fundamentals of working with containers
- Kubernetes in 4 Hours: introduction level course to Kubernetes, very hands on
- CKAD: this course, focus on running applications in Kubernetes
- CKA: Your next course, focus on Kubernetes cluster administration
- Managing Microservices with Containers: a Microservices focussed walk through Kubernetes

# CKAD versus CKA

- CKAD can be considered an entry-level Kubernetes exam
- Candidates will need to create and manage applications in containers using standard ingredients
- CKA is more advanced and measures the candidate's knowledge about managing Kubernetes clusters



# CKAD Crash Course

## 1.1 Understanding Kubernetes Deployment Options

# Understanding Deployment Options

- Kubernetes can run in many environments
  - In a public cloud
  - On-premise in a datacenter
  - Using minikube for testing and developing
- In this course we'll use minikube
- Learn how to set up a Kubernetes cluster using **kubeadm** in my CKA course



# CKAD Crash Course

## 1.2 Using Minikube

# Configuration Requirements

- Minikube offers a complete test environment that runs on Linux, OS-X or Windows
- In this course we'll focus on Minikube as it is easy to setup and has no further dependencies
- You'll also need to have the **kubectl** client on your management platform

# Installing a Lab Environment

- Easy installation on top of Fedora / Ubuntu is provided through a lab setup script
- Clone my git repo: **git clone <https://github.com/sandervanvugt/ckad>**
- Run the **kube-setup.sh** script to setup Minikube as well as the **kubectl** client
- Notice this script is frequently updated
- Step-by-step instructions are in "Setup Guide.pdf" in the Git repo

# Testing Minikube

- **minikube ssh**: logs in to the Minikube host
- **docker ps**: shows all Docker processes on the MK host
- **ps aux | grep localkube**: shows the localkube process on MK host
- **kubectl get all**: shows current resources that have been created



# CKAD Crash Course

## 1.3 Running Your First Application

# Demo

- run an app from minikube dashboard
- use **kubectl create deployment firstnginx --image=nginx** to run another app



# CKAD Crash Course

## 2. Understanding Kubernetes

# What is Kubernetes?

- Kubernetes is "an open-source system for automating deployment, scaling and managing of containerized applications" (<https://kubernetes.io>)
- The word Kubernetes comes from Greek, and means "pilot of the ship"
- A new version is released every 3 months!

# A New Way to do IT

- Kubernetes and Containers offer a new way of working with application workload
- In old IT, applications were installed on a server, and if more capacity was needed, the capacity of the server was scaled up
- Kubernetes offers a platform where the application is offered as containers, connecting to the application is organized through proxies, and scalability is easily organized using Deployment scalability features
- This is the *microservice* approach, where many instances of the server (i.e., the container) are available to respond to a request
- Users are decoupled from the actual container that runs the service they're connecting to. If that container goes down, nobody really cares as services are decoupled

# Dealing with Rapid Changes in Kubernetes

- Don't focus on learning which options to use, focus on understanding and looking up the configuration from the documentation available
- Read the release notes with each new version
- If an item shows as deprecated, stop using it. It will be gone two versions from now



# CKAD Crash Course

## 2.2 Understanding Kubernetes Origins

# Kubernetes Origins

- Kubernetes is based on Google Borg, the internal system that Google has been using for many years to offer access to its applications
  - <https://ai.google/research/pubs/pub43438>
- Borg was also the origin of important Linux kernel features, used in containers nowadays, such as namespaces
- Google donated the Borg specification to the Cloud Native Computing Foundation (CNCF) which was the start of Kubernetes



# CKAD Crash Course

## 2.3 Understanding Kubernetes API Objects

# What is an API?

- An Application Programming Interface (API) is the interface between a client and a server
- An API is used to standardize how to access items provided by the server
- The Kubernetes API defines objects in a Kubernetes environment
- Any command that is used in a Kubernetes environment, is doing somewhat in some way with the API
- Also, all that can be done in Kubernetes, can be explained by exploring API features
- The **kubectl explain** command is an excellent way to explore API features

# Working with Kubernetes

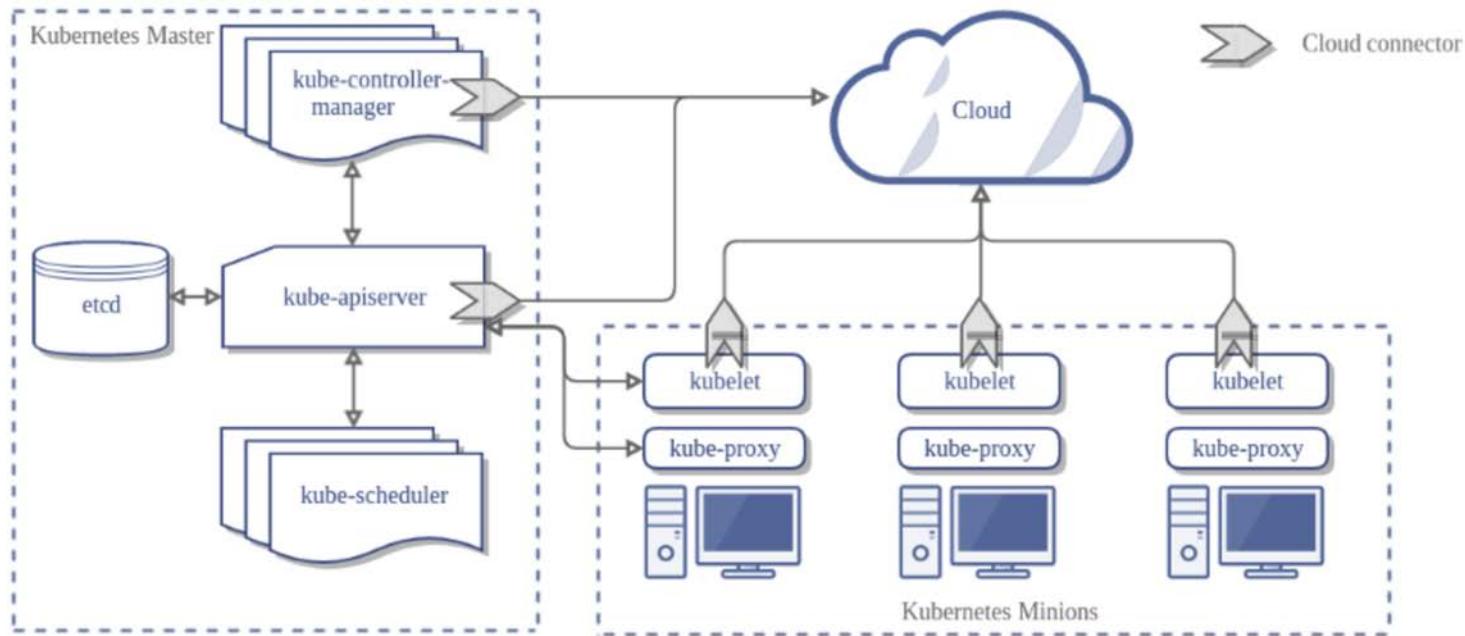
- The **kubectl** command line utility provides convenient administrator access, allowing you to run many tasks against the cluster
- Direct API access allows developers to address the cluster using API calls from custom scripts
- The Kubernetes Console



# CKAD Crash Course

## 2.4 Kubernetes Architecture

# Architecture Overview



# Exam Lab 1

Generate a list of resources provided by the API. Explore the options that are available for creating a service object.



# CKAD Crash Course

## 3. Understanding Managing Options



# CKAD Crash Course

## 4.1 Understanding Main Kubernetes Components

# Kubernetes Main Objects

- Pods: the basic unit in Kubernetes, represents a set of containers that share common resources such as an IP address and persistent storage volumes
- Deployments: standard entity that is rolled out with Kubernetes
- Services: make deployments accessible from the outside by providing a single IP/port combination. Services by default provide access to pods in round-robin fashion using a load balancer
- Persistent Volumes: persistent (networked) storage that can be mounted within a container by using a Persistent Volume Claim



# CKAD Crash Course

## 4.2 Understanding the Kubernetes API

# Understanding the API

- The API defines which objects exist, and how to use these objects
- The main API documentation is here:  
<https://kubernetes.io/docs/reference/generated/kubernetes-api/v1.19/>
- To make extension of the Kubernetes API easier, different groups are used
- Each API group can have its own version number
- See here for more information:  
<https://kubernetes.io/docs/reference/using-api/api-overview/>
- Use **kubectl api-resources** for information about resource types
- Or **kubectl api-versions** for resource and version information

# Understanding API access

- Kubernetes functionality is exposed by the **kube-apiserver** in multiple API groups
- **kubectl** makes secured API requests,
- To make your own API requests using **curl**, the appropriate certificates are required, which is taken care of automatically by **kube-proxy**
- Note that this requires access to CA and certificates, without that you won't see much due to security settings
- Use **kubectl --v=10 get pods** (or any other command) to see the underlying API request

# Connecting to the API

- The API server exposes its functionality through REST
- Assuming that minikube is running, connect to it by running a local proxy first (on your workstation for instance), after which you can connect using curl
  - **kubectl proxy --port=8001&**
  - **curl http://localhost:8001**
- This shows all of the available API paths and groups, providing access to all exposed functions
- Using a proxy allows you to use **curl** and send API calls directly, where the proxy takes care of the appropriate certificates
- Read the documentation for more info about the API groups
- Every API is built according to strict specifications and always contains specific elements

# API Access Example

- On the host that runs kubectl: **kubectl proxy --port=8001 &**
- **curl http://localhost:8001/version**
- **curl http://localhost:8001/api/v1/namespaces/default/pods->** shows the pods
- **curl http://localhost:8001/api/v1/namespaces/default/pods/httpd/**  
shows direct API access to a pod
- **curl -XDELETE**  
**http://localhost:8001/api/v1/namespaces/default/pods/httpd** will delete the httpd pod



# CKAD Crash Course

## 4.3 Using Dashboard

# Using Kubernetes Dashboard

- Kubernetes Dashboard is *not* the recommended way of creating Kubernetes objects
- For an automated way of managing a Kubernetes environment, it is much better to use YAML manifest with the **kubectl** command
- YAML manifests are easy to reproduce, Graphical dashboard interfaces are not



# CKAD Crash Course

## 4.4 Using **kubectl**

# Understanding kubectl

- **kubectl** under the hood uses **curl** to send API requests to the Kubernetes API
- **kubectl** has many subcommands, making it possible to manage all aspects of Kubernetes
- Use **kubectl command --help** for documentation, including examples
- Tip! Also don't forget <https://kubernetes.io/docs>, which is available at the exam also!

# Using `kubectl config`

- The context that `kubectl` uses is stored in `~/.kube/config`
- This context defines which cluster to connect to
- Use `kubectl config view` to view different parts of the current configuration, including current context and namespace

# Exam lab 2

Run a curl command to connect to the Kubernetes API and find out which Pods are created in the Default namespace.



# CKAD Crash Course

## 5. Managing Pods



# CKAD Crash Course

## 5.1 Understanding Pods

# What is a Pod?

- A Pod is an abstraction of a server
  - It runs one or multiple containers within a single name space, exposed by a single IP address
- The Pod is the minimal entity that can be managed by Kubernetes
- Typically, Pods are only started through a Deployment, because "naked" Pods are not rescheduled in case of a node failure

# Managing Pods with kubectl

- Use **kubectl run** to run a *pod* based on a default image
  - **kubectl run ghost --image=ghost:0.9**
  - Notice that **kubectl run** behavior has changed in 1.18
- Use **kubectl** combined with instructions in a YAML file to do anything you'd like
- **kubectl create -f <name>.yaml**
- **kubectl get pods [-o yaml]**
- **kubectl describe pods** shows all details about a pod, including information about containers running within
  - For instance, **kubectl describe pods newhttpd**
- **kubectl edit pod mypod** allows editing of live pods

# Getting Help

- [`https://kubernetes.io/docs`](https://kubernetes.io/docs): also available on the exam!
- **kubectl explain** exposes API resources and fields that are available
- **kubectl explain --recursive** shows all fields



# CKAD Crash Course

## 5.1 Creating a YAML Manifest to Configure Pods

# Using kubectl in a Declarative Way

- The recommended way to work with kubectl, is by writing your manifest files and using **kubectl {create|apply} -f manifest.yaml** to the current objects in your cluster
- This declarative methodology is giving you much more control than the imperative methodology where you create all from the CLI
  - Create new resources: **kubectl create -f nginx.yaml**
  - Push settings from a new manifest: **kubectl replace -f nginx.yaml**
  - Apply settings from a manifest: **kubectl apply -f nginx.yaml**

# Basic YAML Manifest Ingredients

- **apiVersion**: specifies which version of the API to use for this object
- **kind**: indicates the type of object (Deployment, Pod, etc.)
- **metadata**: contains administrative information about the object
- **spec**: contains the specifics for the object
- Use **kubectl explain** to get more information about the basic properties

# Generating YAML Files

- Use a declarative command with command line options to make it easy to generate the YAML code
  - `--dry-run=client`
  - `-o yaml`
- Generate YAML for a running resource
  - `kubectl run busypod --image=busybox --restart=Never --command -- /bin/sh`
  - `kubectl get pods busypod -o yaml > busypod.yaml`
  - Don't forget to cleanup the status section and the **fields** part
- Copy/Paste frome kubernetes.io/docs

# Writing YAML Files

- Indentation identifies parent/child relations
- Use spaces, not tabs for indentation
- Configure **vi** for indentation by adding the following to `~/.vimrc`: **autocmd FileType yaml setlocal ai ts=2 sw=2 et**
- This configuration is already done on the exam environment

# Example YAML File

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
  namespace: default
spec:
  containers:
    - name: busybox
      image: busybox
      command:
        - sleep
        - "3600"
    - name: nginx
      image: nginx
```



# CKAD Crash Course

## 5.3 Understanding Multi-Container Pods

# Understanding Multi-Container Pods

- Typically, just one container is offered through a pod
- Single container Pods are easier to build and maintain
- There are some well-known cases where you might want to run multiple containers in a Pod
  - Sidecar container: a container that enhances the primary application, for instance for logging
  - Ambassador container: a container that represents the primary container to the outside world, such as a proxy
  - Adapter container: used to adopt the traffic or data pattern to match the traffic or data pattern in other applications in the cluster
- When using multi-container pods, the containers typically share data through shared storage

# Understanding the Sidecar Scenario

- A sidecar container is providing additional functionality to the main container, where it makes no sense running this functionality in a separate Pod
- Think of logging, monitoring and syncing
- The essence is that the main container and the sidecar container have access to shared resources to exchange information
- Often, shared volumes (discussed in-depth later) are used for this purpose

# Demo: Sidecar Example

- **kubectl create -f sidecar.yaml**
- **kubectl exec -it sidecar-pod -- curl /bin/bash**
- **yum install -y curl**
- **curl http://localhost/date.txt**

# Understanding the Init Container

- An Init container is a container that runs in a pod to prepare a job for the main container
- The init container will run to completion, after which the primary container can run
- See the exercise **Configure Pod Initialization** in the Kubernetes Documentation (YAML code also available in initpod.yaml)



# CKAD Crash Course

## 5.4 Monitoring Pods

# Monitoring Pods

- Use **kubectl get pods** to get an overview of all pods and their current state
- An event log is separately stored, use **kubectl get events** to print events for the current namespace
- More advanced Pod monitoring should happen from the point of view of application monitoring: don't just monitor the individual pod, but the deployment of which it is a part
- Advisor is a more advanced monitoring agent that integrates with the Kubelet to get performance data about Kubernetes components
- Prometheus / Metric server is a more advanced monitoring solution that allows you to monitor pods as well as nodes

# Exam Lab 3

Start a pod that runs an nginx webserver. Use the name nginx-lab. Find the fastest way possible to do this.



# CKAD Crash Course

## 5.5 Exploring Pod Configuration with **kubectl describe**

# Exploring Pod Configuration

- When deploying a Pod, many parameters are set to a default value
- Use **kubectl describe pod podname-xxx** to see all these parameters and their current setting
- Use documentation at <https://kubernetes.io/docs> for more information about these settings
- Tip! This documentation is available on the exam as well: use it!

# Connecting to a Pod for Further Inspection

- Apart from exploring a Pod externally, you can also connect to it and run commands on the primary container in a pod:
  - **kubectl exec -it nginx-xxx -- sh**
  - From here, run any command to investigate
  - Use Ctrl-p, Ctrl-q to disconnect

# Exam lab 4

Explore the logs of the pod you've created in the previous lab. Which volumes does it have mounted?



# CKAD Crash Course

## 5.6 Exploring Pod Logs

# Using Pod Logs

- Applications running in containers in Pods are not connected to a STDOUT
- Log information is stored in the etcd database, so log information about previous instances is available
- The **kubectl logs** replace the STDOUT
- Use **kubectl logs podname-xxx** to read these logs
- Use **kubectl logs podname-xxx -p** to show logs for previous instances
- Use **kubectl logs podname-xxx --since=1h** to see messages logged in the last hour only



# CKAD Crash Course

## 5.7 Understanding Security Context

# Understanding SecurityContext

- A security context defines privilege and access control settings for a pod or container, and includes the following:
  - Discretionary Access Control which is about permissions to access an object
  - Security Enhanced Linux, where security labels can be applied
  - Running as privileged or unprivileged user
  - Using Linux capabilities
  - AppArmor, which is an alternative to SELinux
  - AllowPrivilegeEscalation which controls if a process can gain more privileges than its parent process
- Notice that **securityContext** exists in pod.spec, as well as pod.spec.containers

# Demo

- **kubectl apply -f securitycontextdemo2.yaml**
- **kubectl exec -it security-context-demo -- sh**
  - ps
  - cd /data; ls -l
  - cd demo; echo hello > testfile
  - ls -l
  - id
- **kubectl explain pods.spec.securityContext** and containers!



# CKAD Crash Course

## 5.8 Managing Jobs and Cron Jobs

# Managing Jobs

- Pods normally are created to run forever
- To create a Pod that runs for a limited duration, use Jobs instead
- Jobs are useful for tasks, like backup, calculation, batch processing and more
- A Pod that is started by a Job must have its restartPolicy set to OnFailure or Never
  - OnFailure will re-run the container on the same Pod
  - Never will re-run the failing container in a new Pod
- A Job is more than just a Pod with the restartPolicy set to never

# Understanding Job Types

3 different job types can be started, which is specified by the **completions** and **parallelism** parameters:

- Non-parallel Jobs: one Pod is started, unless the Pod fails
  - **completions=1**
  - **parallelism=1**
- Parallel Jobs with a fixed completion count: the job is complete after successfully running as many times as specified in jobs.spec.completions
  - **completions=n**
  - **parallelism=m**
- Parallel Jobs with a work queue: multiple jobs are started, when one completes successfully, the job is complete
  - **completions=1**
  - **parallelism=n**

# Demo

- **kubectl create -f simplejob.yaml**
- **kubectl get jobs**
- **kubectl get pods # notice current Pod state**
- **kubectl get pods # notice Pod state after completing the Job command**
- **kubectl get jobs -o yaml**
- **kubectl delete jobs simple-job**
- Edit the Job, and in the spec, add **completions: 3** and run again

# Managing Cron Jobs

- Jobs are used to run a task a specific number of times
- CronJobs are used for tasks that need to run on a regular basis
- When running a Cronjob, a Job will be scheduled
- This job on its turn will start a pod

# Demo

- **kubectl explain CronJob.spec**
- **vim cron-example.yaml**
- **kubectl create -f cron-example.yaml**
- **kubectl get cronjob hello**
- **kubectl get cronjobs**
- **kubectl get jobs --watch (it will take a minute before it runs)**
- **kubectl explain Jobs.spec**
- **kubectl get pods**
- **kubectl get pods hello-nnn -o yaml**
- **kubectl logs hello-nnn**
- **kubectl delete cronjob hello**

# Exam Lab 6

Create a cronjob. It should send the message "hello world" to syslog every 5 minutes.



# CKAD Crash Course

## 5.10 Managing Resource Limitations

# Understanding Resource Limitations

- By default, a Pod will use as much CPU and memory as necessary to do its work
- This can be managed by using Memory/CPU Requests and Limits in `pod.spec.containers.resources`
- Memory as well as CPU limits can be used
- CPU Limits are expressed in millicore or millicpu, 1/1000 of a CPU core
  - So 500 millicore is 0.5 CPU
- Memory limits can be set also, and are converted to the `--memory` option that can be used by the `docker run` command (or anything similar)
- When being scheduled, the `kube-scheduler` ensures that the node running the pods has all requested resources available

# Demo

- **kubectl create -f frontend-resources.yaml**
- **kubectl delete pod frontend # notice this takes some time**
- **kubectl create -f frontend-resources.yaml**
- **kubectl delete pod frontend --grace-period=0 --force # look at the warning. Is this a smart thing to do?**

# Exam Lab 7

Create a Pod that runs an Apache web server. It should not get access to more than 512 MiB of RAM



# CKAD Crash Course

## 5.10 Understanding NameSpaces

# Understanding Namespaces

- At a Linux level namespace implements kernel level resource isolation
- This limitation can be shared in Kubernetes as Kubernetes Namespaces
- Different namespaces can be used to strictly separate between customer resources
- Use **kubectl ... -n namespace** to work in a specific namespace
- Or use **kubectl config set-context --current --namespace=xxx** to set permanently for the current shell (not recommended on the exam!)

# Demo

- `kubectl get all --all-namespaces`
- `kubectl create ns secret`
- `kubectl create -f busybox-ns.yaml`
- `kubectl config set-context --current --namespace=secret`
- `kubectl config view | grep namespace`



# CKAD Crash Course

## 6. Managing Deployments



# CKAD Crash Course

## 6.1 Understanding Deployment Features

# Understanding Deployments

- When running Pods, additional features are needed:
  - Scalability
  - Updates and Update Strategy
  - and more
- That's why you don't run pods, you'll run deployments instead
- When creating an application from Dashboard, you'll create a deployment by default

# Understanding the Deployment Spec

In the deployment spec, specific properties are managed:

- replicas: explains how many copies of each pod should be running
- strategy: explains how Pods should be updated
- selector: uses matchLabels to identify how labels are matched against the Pod
- template: contains the pod specification and is used in a deployment to create Pods

# Demo

- **kubectl create deployment nginxblah --image=nginx**
- **kubectl get deploy nginxblah -o yaml**
- Look for template, selector and matchLabels



# CKAD Crash Course

## 6.2 Managing Deployment Scalability

# Managing Scalability

- Use **kubectl scale deployment my-deployment --replicas=4** to scale the number of currently running replicas
- Alternatively, use **kubectl edit deployment my-deployment** to edit the number of replicas manually

# Understanding ReplicaSets

- ReplicaSets can be used to manage scalability from outside a deployment
- In older versions of Kubernetes, Replicasets were managed as independent resources, in current Kubernetes they should only be managed as part of the deployment!
- From inside a deployment, the spec.replicas parameter is used to indicate the number of desired replicas
- In current Kubernetes, you'll get a ReplicaSet when creating the deployment
- Don't manage scalability through ReplicaSets, the replicas parameter while creating a deployment

# Demo

- show redis-deploy.yaml, run it using **kubectl create -f redis.deploy**
- **kubectl edit deployment redis** → edit the # of replicas
- **kubectl delete rs/redis-nnn** and see how it auto-heals

# Exam Lab 8

Create a deployment. Use the name nginx-lab8, and ensure this deployment runs 3 replicas of an nginx Pod.



# CKAD Crash Course

## 6.3 Managing Deployment History

# Understanding Deployment History

- Major changes cause the deployment to create a new replica set that uses the new properties
- The old replica set is still kept, but the number of Pods will be set to 0
- This makes it easy to undo a change
- **kubectl rollout history** will show the rollout history of a specific deployment, which can easily be reverted as well



# CKAD Crash Course

## 6.4 Understanding Rolling Updates and Rollback

# Understanding Rolling Updates

- It is the task of the deployment to ensure that a sufficient number of Pods is running at all times
- So when making a change, this change is applied as a rolling update: the changed version is deployed and after that has confirmed to be successful, the old version is taken offline
- You can use **kubectl rollout history** to get details about recent transactions
- Use **kubectl rollout undo** to undo a previous change

# Understanding Update Strategy

When a deployment changes, the Pods are immediately updated according to the update strategy:

- Recreate: all Pods are killed and new Pods are created. This will lead to temporary unavailability. Useful if you cannot simultaneously run different versions of an application
- RollingUpdate: updates Pods one at a time to guarantee availability of the application. This is the preferred approach, and you can further tune its behavior

# Using RollingUpdate Options

- maxUnavailable: determines the maximum number of Pods that is upgraded at the same time
- maxSurge: the number of Pods that is allowed to run beyond the desired number of pods as specified in a replica to guarantee minimal availability

# Demo

- **kubectl create -f rolling.yaml**
- **kubectl rollout history deployment**
- **kubectl edit deployment rolling-nginx # change version to 1.15**
- **kubectl rollout history deployment**
- **kubectl describe deployments rolling-nginx**
- **kubectl rollout history deployment rolling-nginx --revision=2**
- **kubectl rollout history deployment rolling-nginx --revision=1**
- **kubectl rollout undo deployment rolling-nginx --to-revision=1**



# CKAD Crash Course

## 6.5 Understanding Labels, Selectors and Annotations

# Understanding Labels

- A label can be used as a key/value pair for further identification of Kubernetes resources
- This can be useful for locating resources at a later stage
- Labels are automatically set when running an app, and the name of the label is based on how the app was started
- Labels are also used by Kubernetes itself for pod selection by deployments and services
  - Deployments are monitoring a sufficient amount of Pods through the run label
  - When creating services using **kubectl expose**, a label is automatically added
- Use the selector option to search for items that have a specific label set
  - **kubectl get pods --selector='run=httpd'**

# Understanding Annotations

- Annotations are used to provide detailed metadata in an object
- Annotations can not be used in queries, they are just to provide additional information
- Think of information about licenses, maintainer and more

# Demo

## MANUALLY SETTING

- **kubectl label deployment ghost state=demo**
- **kubectl get deployments --show-labels**
- **kubectl get deployments --selector state=demo**

## AUTOMATED

- **kubectl create deployment nginx --image=nginx**
- **kubectl describe deployment nginx** → look for label
- **kubectl describe pod nginx-xxx**
- **kubectl label pod nginx-xxx app-** → will remove the auto-assigned run label and start a new pod to meet the requirements
- **kubectl get all --selector app=nginx**

# Exam Lab 9

Find all pods that have been started with the **kubectl run** command

# Exam Lab 10

Create a Cron Job that will run the sleep command for 30 seconds. If after 15 seconds, it still runs, kill it and start the cycle again



# CKAD Crash Course

## 7. Managing Networking



# CKAD Crash Course

## 7.1 Understanding Pod Access Options

# Understanding Pod Access Options

- port forwarding: exposes a port on the local machine that runs the workload
- services: run a kube-proxy as a load balancer on all nodes involved to forward traffic to the appropriate node and uses 4 types:
  - ClusterIP: internal access through a local cluster IP only
  - NodePort: exposes a high port on a fixed external IP
  - LoadBalancer: cloud based endpoint that is externally available to load-balance incoming traffic to the pods
  - externalName: redirects incoming traffic to DNS to further take care of the incoming requests
- ingress: adds scalability: one ingress controller can be used as access point to multiple services



# CKAD Crash Course

## 7.2 Understanding Services

# Understanding Services

- A service is an API object that is used to expose a logical set of Pods
- A service is an abstraction which defines a logical set of Pods and a policy by which to access them
- The set of Pods that is targeted by a service is determined by a selector (which is a label)
- The controller will continuously scan for Pods that match the selector and include these in the service
- This can be useful if communication needs to happen across namespaces

# Understanding Services Decoupling

- Services exist independently from Deployments
- The only thing they do, is watch for a deployment that has a specific label set, based on the Selector that is specified in the service
- That means that one service can provide access to multiple deployments, and while doing so, Kubernetes will automatically load balance between these deployments

# Services and Kube-proxy

- The kube-proxy agent on the nodes watches the Kubernetes API for new services and endpoints
- After creation, it opens random ports and listens for traffic to the clusterIP port, and next redirects traffic to the randomly generated service endpoints

# Understanding Service Types

According to the needs in different environments, different service types are available:

- ClusterIP: the default type. Provides internal access only
- NodePort: allocates a specific node port which needs to be opened on the firewall
- LoadBalancer: currently only implemented in public cloud
- ExternalName: a relatively new object that works on DNS names. Redirection is happening at a DNS level
- Service without selector: use for direct connections based on IP/port, without an endpoint. Useful for connections to database, or between namespaces

# Understanding the ClusterIP Type in Microservices

- The ClusterIP service type doesn't seem to be very useful, as it doesn't allow access to a service from the outside
- For communication of applications in a microservice architecture, it is perfect though: it does expose the application on the cluster IP network, but makes it inaccessible for external users

# Understanding Service Ports

While working with services, different Ports are specified

- targetport: this is the port on the application that the service object connects to
- port: this is what maps from the cluster nodes to the targetport in the application to make the port forwarding connection
- nodeport: this is what is exposed externally



# CKAD Crash Course

## 7.4 Creating Services

# Exposing Deployments

- Use **kubectl expose** to expose a current deployment
- **kubectl expose deployment nginx --port=80 --type=NodePort**
  - Notice this command allocates a random port on all backend nodes, optionally use **targetPort** to define the port that should be used
- **kubectl get svc** will show current services
- **kubectl get svc nginx -o yaml** will show service specifics in YAML
- According to the type that is exposed, you may need to pass more parameters while using **kubectl expose**, see --help for more details

# Demo: Exposing Applications Using Services

- **kubectl create deployment nginxsvc --image=nginx**
- **kubectl scale deployment nginxsvc --replicas=3**
- **kubectl expose deployment nginxsvc --port=80**
- **kubectl describe svc nginxsvc # look for endpoints**
- **kubectl get svc nginx -o=yaml**
- **kubectl get svc**
- **kubectl get endpoints**

# Demo: Accessing Deployments by Services

- **minikube ssh**
- **curl <http://svc-ip-address>**
- **exit**
- **kubectl edit svc nginxsvc**  
...  
**protocol: TCP**  
**nodePort: 32000**  
**type: NodePort**
- **kubectl get svc**
- (from host): **curl http://\$(minikube ip):32000**



# CKAD Crash Course

## 7.5 Managing Service Manifest Files

# Demo

- see service.yml
- Note that the essence is in the selector app: MyApp which will make this configuration applies to all objects that meet this selector
- Notice this is a clusterIP type, this clusterIP is used by the service proxies



# CKAD Crash Course

## 7.6 Understanding Services and DNS

# Understanding Services and DNS

- With services exposing themselves on dynamic ports, resolving service names can be challenging
- As a solution, a DNS service is included by default in Kubernetes and this DNS service is updated every time a new service is added
- So DNS name lookup from within one pod to any exposed service happens automatically

# Demo

Demo:

1. Create busybox.yaml

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: busybox2
```

```
  namespace: default
```

```
spec:
```

```
  containers:
```

```
    - image: busybox
```

```
      name: busy
```

```
      command:
```

```
        - sleep
```

```
        - "3600"
```

2. Create the pod, using **kubectl create -f busybox.yaml**

3. Use **kubectl get svc** to validate the name of any exposed service

4. Use **kubectl exec -ti busybox2 -- nslookup http**

5. You'll see the IP address being resolved, thus providing proof that DNS is working correctly

# Exam lab 12

- Configure a Service for the Nginx deployment you've created earlier. Ensure that this service makes Nginx accessible through port 80, using the ClusterIP type. Verify that works
- After making the service accessible this way, change the type to NodePort and expose the service on port 32000
- Verify the service is accessible on this node port



# CKAD Crash Course

## 8. Managing Ingress



# CKAD Crash Course

## 8.1 Understanding Ingress

# Understanding Ingress

- Ingress exposes HTTP and HTTPS routes from outside the cluster to services within the cluster
- Other traffic should be handled using service objects
- Traffic routing is controlled by rules defined on the Ingress resource
- Ingress can be configured to do the following
  - Give Services externally-reachable URLs
  - Load balance traffic
  - Terminate SSL/TLS
  - Offer name based virtual hosting
- An Ingress controller is used to realize the Ingress
  - On Minikube, use **minikube addons enable ingress**

# Understanding Ingress Controllers

- The Ingress controller is a network specific solution that programs an external load balancer
- Many Ingress controllers exist:
  - nginx: <https://kubernetes.github.io/ingress-nginx/> (minikube default)
  - haproxy: <https://www.haproxy.com/blog/dissecting-the-haproxy-kubernetes-ingress-controller/>
  - traefik: <https://docs.traefik.io>
  - kong: <https://konghq.com/solutions/kubernetes-ingress/>
  - contour: <https://octetz.com/posts/contour-adv-ing-and-delegation>
- To enable the minikube addon, use **minikube addons enable ingress**



# CKAD Crash Course

## 8.2 Configuring Ingress

# Sample Ingress Manifest

```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: test-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - http:
    paths:
    - path: /testpath
      backend:
        serviceName: test
        servicePort: 80
```

# Demo: Creating Ingress - 1

Note: this demo continues on the demo in lesson 7.4

- **minikube addons enable ingress**
- **kubectl get deployment**
- **kubectl get svc nginxsvc**
- **curl http://\$(minikube ip):32000**
- **vim nginxsvc-ingress.yaml**
- **kubectl apply -f nginxsvc-ingress.yaml**
- **kubectl get ingress** - wait until it shows an IP address, this takes 3 minutes
- **sudo vim /etc/hosts**
  - **\$(minikube ip) nginxsvc.info**
- **curl nginxsvc.info**

# Demo: Creating Ingress - 2

- **kubectl create deployment newdep --image=gcr.io/google-samples/hello-app:2.0**
- **kubectl expose deployment newdep --port=8080**
- Add the following to **nginxsvc-ingress.yaml**

```
- path: /hello
  pathType: Prefix
  backend:
    service:
      name: newdep
      port:
        number: 8080
```



# CKAD Crash Course

## 9. Managing Storage



# CKAD Crash Course

## 9.1 Understanding Storage Options

# Understanding Storage Options

- Files stored in a container will only live as long as the container itself
- Pod Volumes can be used to allocated storage that outlives a container and stays available during pod lifetime
  - Pod Volumes can refer to internal as well as external storage
- A persistent volume allows administrators to decouple storage from the actual Pod
- To use persistent volumes, persistent volume claims (PVC) are used to request access to specific storage
  - PVC allow developers to request storage without knowing about storage specifics
- ConfigMaps are specific objects that connect to configuration files
- Secrets do the same, but by encoding the data they contain



# CKAD Crash Course

## 9.2 Configuring Volume Storage

# Configuring Volumes

- First, decide if you want to use a regular volume or a persistent volume
- To create a regular volume, the Pod needs to define the volume in `spec.volumes`.
- Next, the container mounts that volume in `spec.containers.volumemounts`
- For use of persistent volume, additional external objects are needed

# Demo

1. **kubectl create -f morevolumes.yaml**
2. **kubectl get pods morevol2**
3. **kubectl describe pods morevol2 | less** ## verify there are two containers in the pod
4. **kubectl exec -ti morevol2 -c centos1 -- touch /centos1/test**
5. **kubectl exec -ti morevol2 -c centos2 -- ls -l /centos2**



# CKAD Crash Course

## 9.3 Configuring PV Storage

# Understanding Persistent Volumes

- A Persistent Volume is a storage abstraction that is used to store persistent data
  - Use **persistentVolume** to define it
  - Different types (NFS, iSCSI, CephFS and many more) are available
- Using claims with Persistent Volumes creates portable Kubernetes storage that allow you to use volumes, regardless of the specific storage provider
  - Use **persistentVolumeClaims**
  - The persistent volume claim talks to the available backend storage provider and dynamically uses volumes that are available on that storage type
- Kubernetes will use Persistent Volumes according to the availability of the requested volume accessModes and capacity

# Understanding StorageClass

- In a decoupled environment, persistent storage is taken care of by persistent volumes
- In cluster environments, a StorageClass can be used
- StorageClass is default cluster storage that replaces the need to manually configure persistent volumes at all times



# CKAD Crash Course

## 9.4 Configuring PVCs

# Understanding Persistent Volume Claims

- The Persistent Volume Claims requests access to storage provided by Persistent Volumes according to specific properties
  - accessModes
  - availability of resources
- The persistent volume sets a name to the PVC, which can be used in the Pod which is accessing the storage



# CKAD Crash Course

## 9.5 Configuring Pod Storage

# Exam Lab 13: Setting up Storage

- Configure a PV that is using HostPath storage
- Set up the storage accessMode such that only multiple pods may access the storage at the same time
- Configure a Pod that runs an httpd web server to mount the default web server documentroot /var/www/html on the HostPath storage



# CKAD Crash Course

## 10. Managing ConfigMaps and Secrets



# CKAD Crash Course

## 10.1 Understanding ConfigMaps

# Understanding ConfigMap

- ConfigMaps can be used to separate dynamic data from static data in a Pod
- They are not encoded or encrypted
- They can be used in three different ways:
  - Make variables available within a Pod
  - Provide command line arguments
  - Mount them on the location where the application expects to find a configuration file
- Secrets are encoded ConfigMaps which can be used to store sensitive data
- ConfigMaps must be created before the pods that are using them

# Understanding ConfigMap Sources

- ConfigMaps can be created from different sources
  - Directories: uses multiple files in a directory
  - Files: puts the contents of a file in the ConfigMap
  - Literal Values: useful to provide variables and command arguments that are to be used by a Pod



# CKAD Crash Course

## 10.2 Creating ConfigMaps

# Procedure Overview

- Start by defining the ConfigMap and create it
  - Consider the different sources that can be used for ConfigMaps
  - **kubectl create cm myconf --from-file=my.conf**
  - **kubectl create cm variables --from-env-file=variables**
  - **kubectl create cm special --from-literal=VAR3=cow --from-literal=VAR4=goat**
  - Verify creation, using **kubectl describe cm <cmname>**
- Use **--from-file** to put the contents of a config file in the configmap
- Use **--from-env-file** to define variables
- Use **--from-literal** to define variables or command line arguments

# Procedure Overview - 2

- To include *variables* from a ConfigMap in a pod:

`envFrom:`

- `configMapRef:`

  - `name: ConfigMapName`

- To include *config files* from a ConfigMap in a pod:

`volumes:`

- `configMap:`

  - `name: ConfigMapName`

  - `items:`

  - `key: my-custom.conf`

  - `path: default.conf`

# Demo: Creating a ConfigMap from a File

- Show contents of variables (in github)
- Create the ConfigMap: **kubectl create cm variables --from-env-file=variables**
- Verify creation: **kubectl describe cm variables**
- Create a Pod: **kubectl create -f cm-test-pod.yml**
- Check that the variables are available: **kubectl logs po/test** (or whatever the name is)

# Demo: Configuring a ConfigMap from a Literal

- **kubectl create cm morevars --from-literal=VAR3=goat --from-literal=VAR4=cow**
- **kubectl get cm/morevars**

# Demo: Using ConfigMaps for ConfigFiles

- Create the ConfigMap: **kubectl create cm nginx-cm --from-file nginx-custom-config.conf**
- Check the contents of the ConfigMap: **kubectl get configmap/nginx-cm -o yaml**
- Next, create the Pod: **kubectl create -f nginx-cm.yml**
- Check the config file: **kubectl exec -it nginx-cm /bin/bash**
- **cat /etc/nginx/conf.d/default.conf**



# CKAD Crash Course

## 10.3 Understanding Secrets

# Understanding Secrets

- Secrets allow for storage of sensitive data such as passwords, Auth tokens and SSH keys
- Using secrets makes that the data doesn't have to be put in a Pod, and reduces the risk of accidental exposure
- Some secrets are automatically created by the system, users can also use secrets
- System created secrets are important for Kubernetes resources to connect to other cluster resources
- Secrets are used by Pods as files in a volume, or used by kubelet when pulling images for the pod

# Understanding Secret Types

- Three types of secret are offered
  - docker-registry: used for connecting to a Docker registry
  - TLS: creates a TLS secret
  - generic: creates a secret from a local file, directory or literal value

# Understanding Built-in Secrets

- Kubernetes automatically creates secret that contain credentials for accessing the API, and automatically modifies the pods to use this type of secret
- Secrets are not encrypted, they are encoded
- Use **kubectl describe pods <podname>** and look for the mount section to see it



# CKAD Crash Course

## 10.4 Creating Secrets

# Creating Your Own Secrets

- Creating secrets is very similar to creating ConfigMaps
- Use **kubectl create secret ... --from-file=...** to create it from a file
- Use **kubectl create secret ... --from-literal=...** to create it from a string provided on the command line

# Demo

- **kubectl create secret generic my-secret --from-file=ssh-privatekey=/home/student/.ssh/id\_rsa --from-literal=passphrase=password**
- **kubectl describe secrets my-secret**

# Creating Secrets from YAML Files

- When creating Secrets from YAML files, they must first be encoded using the **base64** command
- Note that *encoded* is not the same as *encrypted*, users can easily decode using **echo <string> | base64 -d**

# Demo

This demo shows how to manually create secrets, by base64 encoding them and putting them in a yaml file. Alternatively, use **kubectl create secret**

- `echo -n 'lisa' | base64 >> secret-yaml.yaml`
- `echo -n 'password' | base64 >> secret-yaml.yaml`
- `cat secret-yaml.yaml`
- `kubectl create -f secret-yaml.yaml`
- `kubectl explain secret # notice the differences between data: and stringData:`



# CKAD Crash Course

## 10.5 Configuring Pods to Use Secrets

# Using Secrets

- Secrets can be used by Pods in two ways:
  - As environment variables
  - Mounted as volumes

# Demo: mounting secrets as volumes

- **kubectl create secret generic secretstuff --from-literal=password=password --from-literal=user=linda**
- **kubectl create -f pod-secret.yaml**
- **kubectl describe secretbox2**
- **kubectl exec -ti secretbox2 /bin/sh**
- **cat /secretstuff/password**

# Demo: secrets as variables

- **kubectl create secret generic mysql --from-literal=password=root**
- **kubectl get secret mysql -o yaml**
- continue with pod-secret-as-var.yaml
- show how to get there, using **kubectl explain pod.spec.containers.env.valueFrom**
- **kubectl create -f pod-secret-as-var.yaml**
- **kubectl exec -ti mymysql /bin/bash**
- **env**



# CKAD Crash Course

## 10. Managing Service Accounts

# Understanding Service Accounts

- All actions in a Kubernetes Cluster need to be authenticated and authorized
- Service Accounts are used for basic authentication from within the Kubernetes cluster
- Roles and role bindings are used for authorization
- Every Pod uses the Default ServiceAccount to contact the API server
- This default service account allows a resource to get information from the API server, but not much else
- Each service account uses a secret to automount API credentials

# Understanding ServiceAccount Secrets

- Service Accounts come with a secret
- This secret contains API credentials
- By specifying the service account to be used by a pod, the service account secret is auto-mounted to provide API access credentials

# Understanding Service Account Permissions

- A Cluster Administrator can configure RBAC authorization to specify what exactly can be accessed by a service account
- Custom service accounts can be created, and configured with additional permissions
- For additional permissions, namespace or cluster scope roles and rolebindings are used

# Managing Service Accounts

- Every namespace has a default service account called default
- Additional service accounts can be created to provide additional access to resources
- After creating the additional service account, authorization needs to be set up

# Creating Additional Service Accounts

- Service accounts are easily created using declarative or imperative way
- While creating a service account, a secret is automatically created to have the service account connect to the API
- This secret is mounted in Pods using the service account
- Authorization plugins can be used to set permissions on service accounts
- **kubectl create serviceaccount mysa**

# Demo: Using Service Accounts - 1

1. Create a pod, using the standard service account: **kubectl apply -f mypod.yaml**
2. Use **kubectl get pods mypod -o yaml** to check current SA configuration
3. Access the pod using **kubectl exec -it mypod -- sh**, try to list pods using CURL on the API:
  1. **apk add --update curl**
  2. **curl https://kubernetes/api/v1 --insecure** will be forbidden
4. Use the default service account token and try again:
  1. **TOKEN=\$(cat /run/secrets/kubernetes.io/serviceaccount/token)**
  2. **curl -H "Authorization: Bearer \$TOKEN" https://kubernetes/api/v1/ --insecure**
5. Try the same, but this time to list pods - it will fail:
  1. **curl -H "Authorization: Bearer \$TOKEN"  
https://kubernetes/api/v1/namespaces/default/pods/ --insecure**

# Demo: Using Service Accounts - 2

1. Create a service account: **kubectl apply -f mysa.yaml**
2. Define a role that allows to list all pods in the default namespace: **kubectl apply -f list-pods.yaml**
3. Define a role binding that binds the mysa to the role just created: **kubectl apply -f list-pods-mysa-binding.yaml**
4. Create a Pod that uses the mysa SA to access this role: **kubectl apply -f mysapod.yaml**
5. Use the mysa service account token and try again:
  1. **TOKEN=\$(cat /run/secrets/kubernetes.io/serviceaccount/token)**
  2. **curl -H "Authorization: Bearer \$TOKEN" https://kubernetes/api/v1/ --insecure**
6. Try the same, but this time to list pods:
  1. **curl -H "Authorization: Bearer \$TOKEN"  
https://kubernetes/api/v1/namespaces/default/pods/ --insecure**



# CKAD Crash Course

## 11. Troubleshooting



# CKAD Crash Course

## 11.1 Generic Tips

# Troubleshooting Procedure

- Verify syntax before adding anything
- Use standard Linux tools for troubleshooting
  - Consider starting busybox if tools are absent
- Use **kubectl describe** to see what happens
- Use **kubectl logs** to see what a container in a Pod is doing
- Use **kubectl get events** to get the generic event log
- Consider limitations: logging activity across all nodes is not part of Kubernetes and provided by external projects like Fluentd or Prometheus
- In this lesson we'll focus on **kubectl** commands that can be used

# Verifying Syntax

- Without Internet Access: use **kubectl explain**
  - **kubectl explain** explains API objects
  - To get information about a specific part of the object, use dotted notation:  
**kubectl explain pod.spec**
- With Internet Access: use <https://kubernetes.io/docs>
  - Use the Search feature for examples
  - Most of the useful examples are in the Tasks section of the Docs
  - You are allowed to access kubernetes.io/docs during the exam

# Demo

- This demo is mainly about explain and log
- Clean up previously used cm objects and the test1 pod that was created from cm-test-pod.
- edit **cm-test-pod.yaml** and remove restartPolicy: Never
- Use **kubectl create cm variables --from-file=variables**
- Use **kubectl create -f cm-test-pod1.yaml**
- Use **kubectl get pods**: you'll see a Crashloopback failure
- Use **kubectl describe pods test1** and scroll through the output
- Use **kubectl logs test1** to see that the VAR1 and VAR2 are in the list of the logs – which basically means it was successful
- Use **kubectl explain pod.spec**, scroll through options and look for restartPolicy
- Apply restartPolicy to fix



# CKAD Crash Course

## 11.2 Using Probes

# Understanding Probes

- Probes can be used to test access to Pods
- A readinessProbe is used to make sure a Pod is not published as available until the readinessProbe has been able to access it
- The livenessProbe is used to continue checking the availability of a Pod
- A probe is a simple test that verifies that the application is reacting

# Understanding probe Types

- command: a command is executed and returns a zero exit value
- HTTP request: an HTTP request returns a response code between 200 and 399
- TCP socket: connectivity to a TCP socket (available port) is successful

# Demo

- **kubectl create -f busybox-ready.yaml**
- **kubectl get pods** note the READY state, which is set to 0/1, which means that the pod has successfully started, but is not considered ready.
- **kubectl edit pods busybox-ready** and change /tmp/nothing to /etc/hosts. Notice this is not allowed.
- **kubectl exec -it busybox-ready -- /bin/sh**
  - `touch /tmp/nothing; exit`
- **kubectl get pods** at this point we have a pod that is started

# Demo 2

- **kubectl create -f nginx-probes.yaml**
- **kubectl get pods**



# CKAD Crash Course

What's next?

# Follow up courses

- Certified Kubernetes Administrator (CKA) Crash Course
- Building Microservices with Containers